

Experiment 6 – SVR Write-up

Name: Ashish Nanda

Roll No.: J041

```
sklearn.svm.SVR(*, kernel='rbf', degree=3, gamma='scale', coef0=0.0,
tol=0.001, C=1.0, epsilon=0.1, shrinking=True, cache_size=200,
verbose=False, max_iter=- 1)
```

Parameters

kernel{'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'}, default='rbf'

Specifies the kernel type to be used in the algorithm. It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or a callable. If none is given, 'rbf' will be used. If a callable is given it is used to precompute the kernel matrix.

degreeint, default=3

Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.

gamma{'scale', 'auto'} or float, default='scale'

Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.

- if gamma='scale' (default) is passed then it uses $1 / (n_features * X.var())$ as value of gamma,
- if 'auto', uses $1 / n_features$.

coef0float, default=0.0

Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.

tolfloat, default=1e-3

Tolerance for stopping criterion.

Cfloat, default=1.0

Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l2 penalty.

epsilonfloat, default=0.1

Epsilon in the epsilon-SVR model. It specifies the epsilon-tube within which no penalty is associated in the training loss function with points predicted within a distance epsilon from the actual value.

shrinkingbool, default=True

Whether to use the shrinking heuristic.

cache_sizefloat, default=200

Specify the size of the kernel cache (in MB).

verbosebool, default=False

Enable verbose output. Note that this setting takes advantage of a per-process runtime setting in libsvm that, if enabled, may not work properly in a multithreaded context.

max_iterint, default=-1

Hard limit on iterations within solver, or -1 for no limit.

Attributes

class_weight_ndarray of shape (n_classes,)

Multipliers of parameter C for each class. Computed based on the class_weight parameter.

coef_ndarray of shape (1, n_features)

Weights assigned to the features (coefficients in the primal problem). This is only available in the case of a linear kernel.

`coef_` is readonly property derived from `dual_coef_` and `support_vectors_`.

`dual_coef_` *ndarray of shape (1, n_SV)*

Coefficients of the support vector in the decision function.

`fit_status_` *int*

0 if correctly fitted, 1 otherwise (will raise warning)

`intercept_` *ndarray of shape (1,)*

Constants in decision function.

`n_support_` *ndarray of shape (n_classes,), dtype=int32*

Number of support vectors for each class.

`shape_fit_` *tuple of int of shape (n_dimensions_of_X,)*

Array dimensions of training vector X.

`support_` *ndarray of shape (n_SV,)*

Indices of support vectors.

`support_vectors_` *ndarray of shape (n_SV, n_features)*

Support vectors.

Working of Support Vector Regression (SVR)

Support Vector Regression is a supervised learning algorithm that is used to predict discrete values. The basic idea behind SVR is to find the best fit line, it is the hyperplane that has the maximum number of points

Unlike other Regression models that try to minimize the error between the real and predicted value, the SVR tries to fit the best line within a threshold value.

The threshold value is the distance between the hyperplane and boundary line.

The fit time complexity of SVR is more than quadratic with the number of samples which makes it hard to scale to datasets with more than a couple of 10000 samples.

The model produced by Support Vector Regression depends only on a subset of the training data, because the cost function ignores samples whose prediction is close to their target.

For large datasets, Linear SVR or SGD Regressor is used. Linear SVR provides a faster implementation than SVR but only considers the linear kernel.